



Deep Learning for Image Understanding

Challenge for Deep Learning




- Example: 200x200 image
 - Fully connected NN, 40,000 hidden units=1600000000 parameters to link the input and layer 1.
 - Locally connected, 40,000 hidden units 10x10 fields=4000000 parameters
 - Local connections capture local dependencies.



Shared weights & Convolution

- Features that are useful on one part of the image are probably useful elsewhere.
 - Sparse representation
 - “Image is sparse”- A. Leven
- Shift equivalent processing (spatial property)
 - When the input shifts, the outputs also shifts.
- Convolution
 - With a learned kernel (or filter): $A_{ij} = \sum_{kl} W_{kl} X_{i+j, k+l}$
 - The filtered “image” is called a feature map.


- Example: 200x200 image
 - 10 filters of size 10x10
 - 10 feature maps of size 200x200
 - 400,000 hidden units with 10x10 fields=1000 parameters



Computation workload is reduced!


Why 10 filters?

- Detect multiple motifs at each location.
- The collection of units looking at the same patch is akin to a feature vector for that patch.
- The result is a 3D array, where each slice is a feature map.





Early hierarchical feature model for vision

- [Hubel & Wiesel 1962]
 - Simple cells detect local features.
 - Complex cells “pool” the outputs of simple cells within a retinotopic neighborhood.




Pooling

- Pooling
 - Spatial pooling
 - Non-overlapping/overlapping regions
 - Average or max
 - Usually combine with subsampling
 - Role of spatial pooling
 - Invariance to small transformation
 - Larger receptive field
 - Smoothness
 - Reduce variants



Max




Averag

e

Pooling


■ Pooling

- Pooling across feature maps
 - Additional form of inter-feature competition
 - Max-pooling
- Role of pooling across maps
 - Find distinct features
 - Reduce variants



A development: cuda-convnet

- Though convolution and pooling/subsampling reduce the variants greatly, the training of CNN is still time consuming.
- Hence, cuda-convnet is presented by extending LeNet [Lecun, 98, 06]
 - A. Krizhevsky, I. Sutskever, G. E. Hinton, ImageNet classification with deep convolutional neural networks, NIPS 2012.
 - Code: <https://code.google.com/p/cuda-convnet/>




Cuda-convnet

■ Features


- Trained with gradient decent on two nVidia GPUs for about a week @ ImageNet
- 650,000 neurons
- 60,000,000 parameters
- 7 hidden “weight” layers
- Final feature layer: 4096-dimensional

○ Convolutional layer: convolves its input with a bank of 3D filters, then applies “sigmoid”-like operation.


□ Fully-connected layer: applies linear filters to its input, then applies “sigmoid”-like operation.



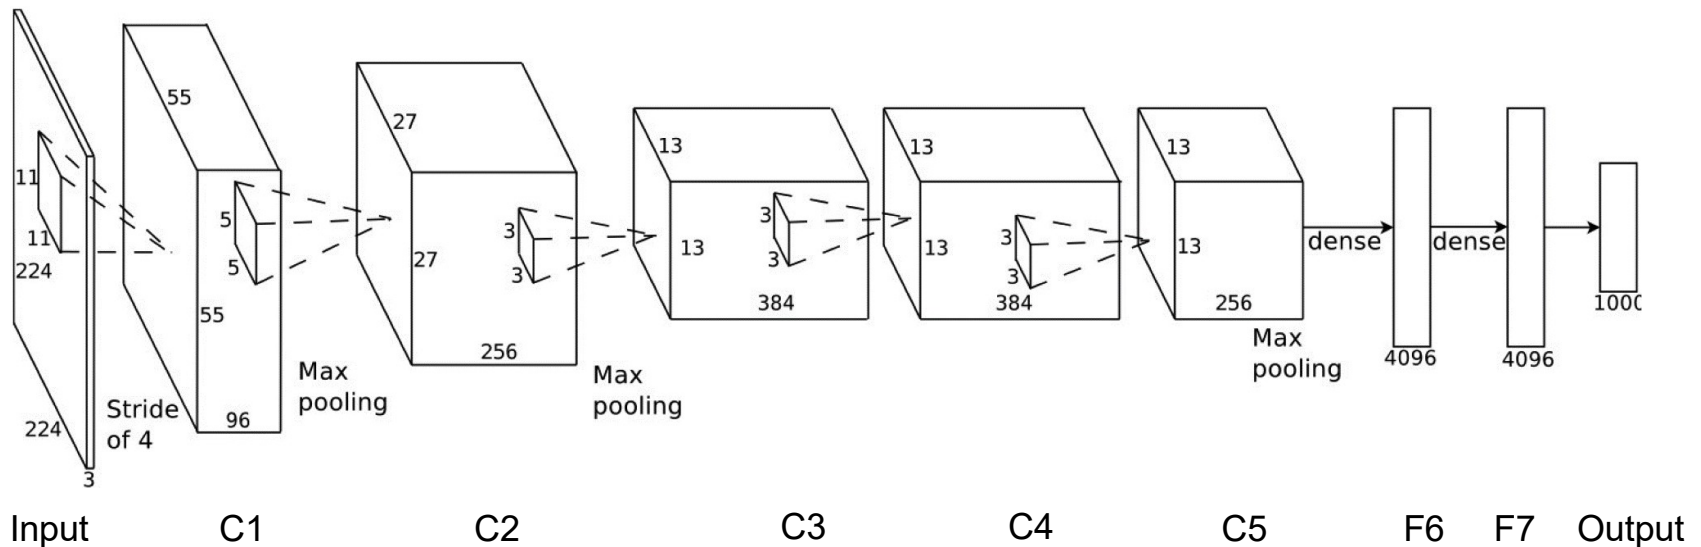
- Input: $3@224 \times 224 \rightarrow C1: 96@55 \times 55$
 - 96 $11 \times 11 \times 3$ convolutional kernels (3D filters) are applied on $224 \times 224 \times 3$ input image.
 - Sliding with the stride of 4.
 - 96 55×55 feature volumes are obtained.




- $C1: 96@55 \times 55 \rightarrow C2: 256@27 \times 27$
 - Max pooling: choose the maximal value in each 2×2 neighborhood, change C1 to be $96@27 \times 27$.
 - 256 $5 \times 5 \times 96$ 3D kernels are applied on C1: $96@27 \times 27$.
 - 256 27×27 feature maps are obtained.




- C2: $256@27 \times 27 \rightarrow$ C3: $384@13 \times 13$
 - Max-pooling: choose the maximal value in each 2×2 neighborhood, change C2 to: $256@13 \times 13$
 - 384 $3 \times 3 \times 256$ 3D kernels are applied on C2: $256@13 \times 13$.
 - 384 13×13 feature maps are obtained.




- C3: $384@13 \times 13 \rightarrow$ C4: $384@13 \times 13$
 - 384 $3 \times 3 \times 384$ 3D kernels are applied on C3.
 - No intervening pooling or normalization.
- C4: $384@13 \times 13 \rightarrow$ C5: $256@13 \times 13$
 - 256 $3 \times 3 \times 384$ kernels are applied on C4.
 - No intervening pooling or normalization.



- C5: 256@13x13 → F6: 4096@1x1
 - Max pooling across 256 different feature maps.
 - Fully connected neural network between the pooled C5 and F6.



- F6: 4096@1x1 \rightarrow F7: 4096@1x1
 - Fully connected neural network between F6 and F7.
- F7: 4096@1x1 \rightarrow Output: 1000@1x1
 - Fully connected neural network between F7 and output layer.



Tasks for which CNNs are the best

- Handwriting recognition: MNIST (many), Arabic HWX (IDSIA)
- OCR in the wild [2011]: StreetView House Numbers (NYU)
- Traffic sign recognition [2011]: GTSRB competition (IDSIA, NYU)
- Pedestrian detection [2013]: INRIA datasets and others (NYU)
- Volumetric brain image segmentation [2009]: Connectomics (IDSIA, MIT)
- Human action recognition [2011]: Hollywood II dataset (Stanford)
- Object recognition [2012]: ImageNet competition.
- Scene Parsing [2012]: Stanford bgd, SiftFlow, Barcelona (NYU)
- Scene parsing from depth images [2013]: NYU RGB-D dataset (NYU)
- Speech recognition [2012]: Acoustic modeling (IBM and Google)
- Breast cancer cell mitosis detection [2011]: MITOS (IDSIA)

Some state-of-the-art performance

- Traffic sign recognition
 - German Traffic Sign Reco Bench
 - 99.2% accuracy




□ House Number Recognition

- Street View House Numbers
- 94.3% accuracy



Scene parsing/labeling on RGB-D images

■ wall	■ books	■ chair	■ furniture	■ sofa	■ object	■ TV
■ bed	■ ceiling	■ floor	■ pict./deco	■ table	■ window	■ unknow



Ground truths

Our results









Semantic segmentation on RGB-D images



Result

Ground truth

Validation classification

			
mite	container ship	motor scooter	leopard
<div></div> <div>mite</div> <div>black widow</div> <div>cockroach</div> <div>tick</div> <div>starfish</div>	<div></div> <div>container ship</div> <div>lifeboat</div> <div>amphibian</div> <div>fireboat</div> <div>drilling platform</div>	<div></div> <div>motor scooter</div> <div>go-kart</div> <div>moped</div> <div>bumper car</div> <div>golfcart</div>	<div></div> <div>leopard</div> <div>jaguar</div> <div>cheetah</div> <div>snow leopard</div> <div>Egyptian cat</div>
			
grille	mushroom	cherry	Madagascar cat
<div></div> <div>convertible</div> <div>grille</div> <div>pickup</div> <div>beach wagon</div> <div>fire engine</div>	<div></div> <div>agaric</div> <div>mushroom</div> <div>jelly fungus</div> <div>gill fungus</div> <div>dead-man's-fingers</div>	<div></div> <div>dalmatian</div> <div>grape</div> <div>elderberry</div> <div>ffordshire bullterrier</div> <div>currant</div>	<div></div> <div>squirrel monkey</div> <div>spider monkey</div> <div>titi</div> <div>indri</div> <div>howler monkey</div>

Validation localization




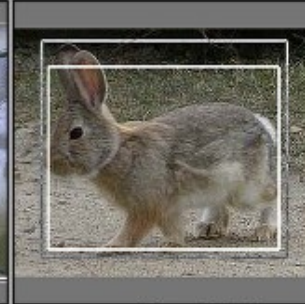




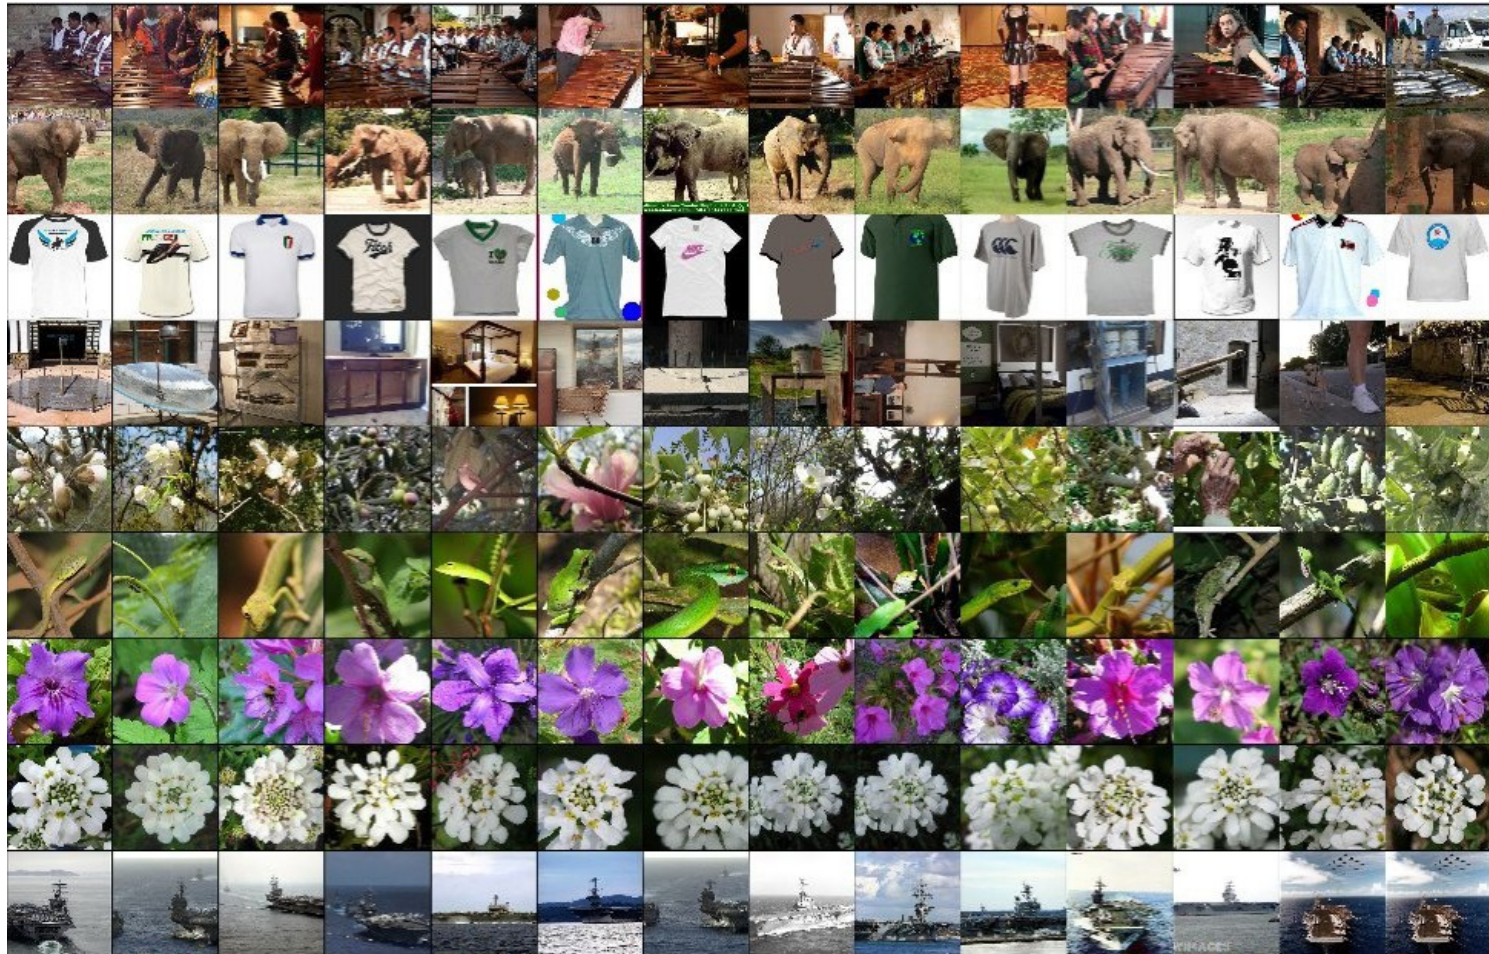
			
bookshop	coyote	cradle	wood rabbit
<div>balance beam</div> <div>cinema</div> <div>marimba</div> <div>parallel bars</div> <div>computer keyboard</div>	<div>grey fox</div> <div>kit fox</div> <div>red fox</div> <div>coyote</div> <div>dhole</div>	<div>cradle</div> <div>bassinet</div> <div>diaper</div> <div>crib</div> <div>bath towel</div>	<div>hare</div> <div>wood rabbit</div> <div>grey fox</div> <div>coyote</div> <div>wallaby</div>
			
bottlecap	harvester	garter snake	Walker hound
<div>bottlecap</div> <div>magnetic compass</div> <div>puck</div> <div>stopwatch</div> <div>disk brake</div>	<div>harvester</div> <div>thresher</div> <div>plow</div> <div>tractor</div> <div>tow truck</div>	<div>diamondback</div> <div>leatherback turtle</div> <div>sandbar</div> <div>echidna</div> <div>armadillo</div>	<div>beagle</div> <div>Walker hound</div> <div>English foxhound</div> <div>muzzle</div> <div>Italian greyhound</div>

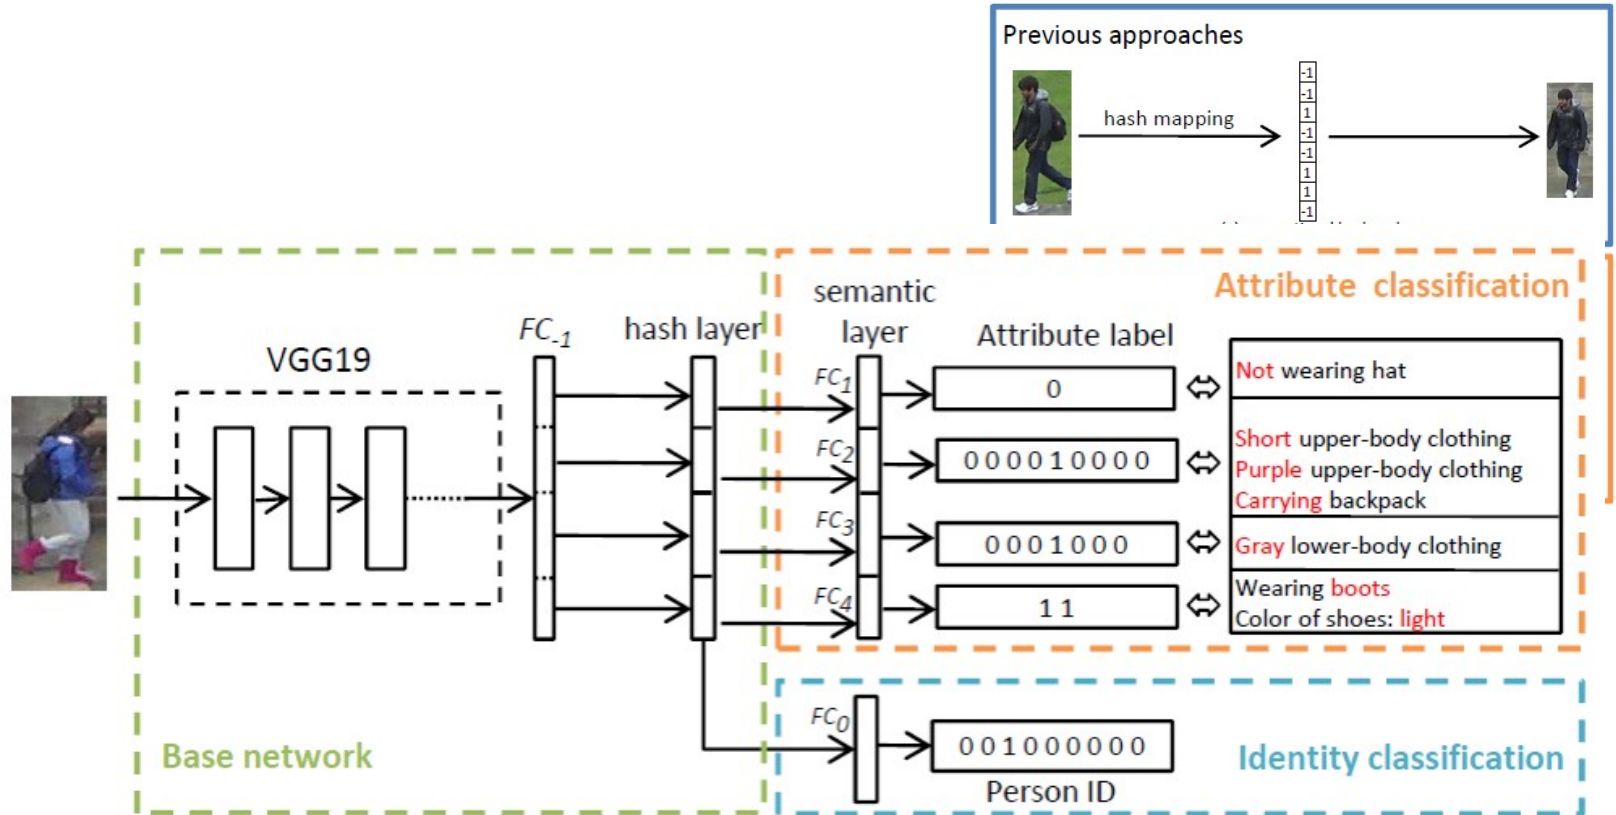
Image retrieval



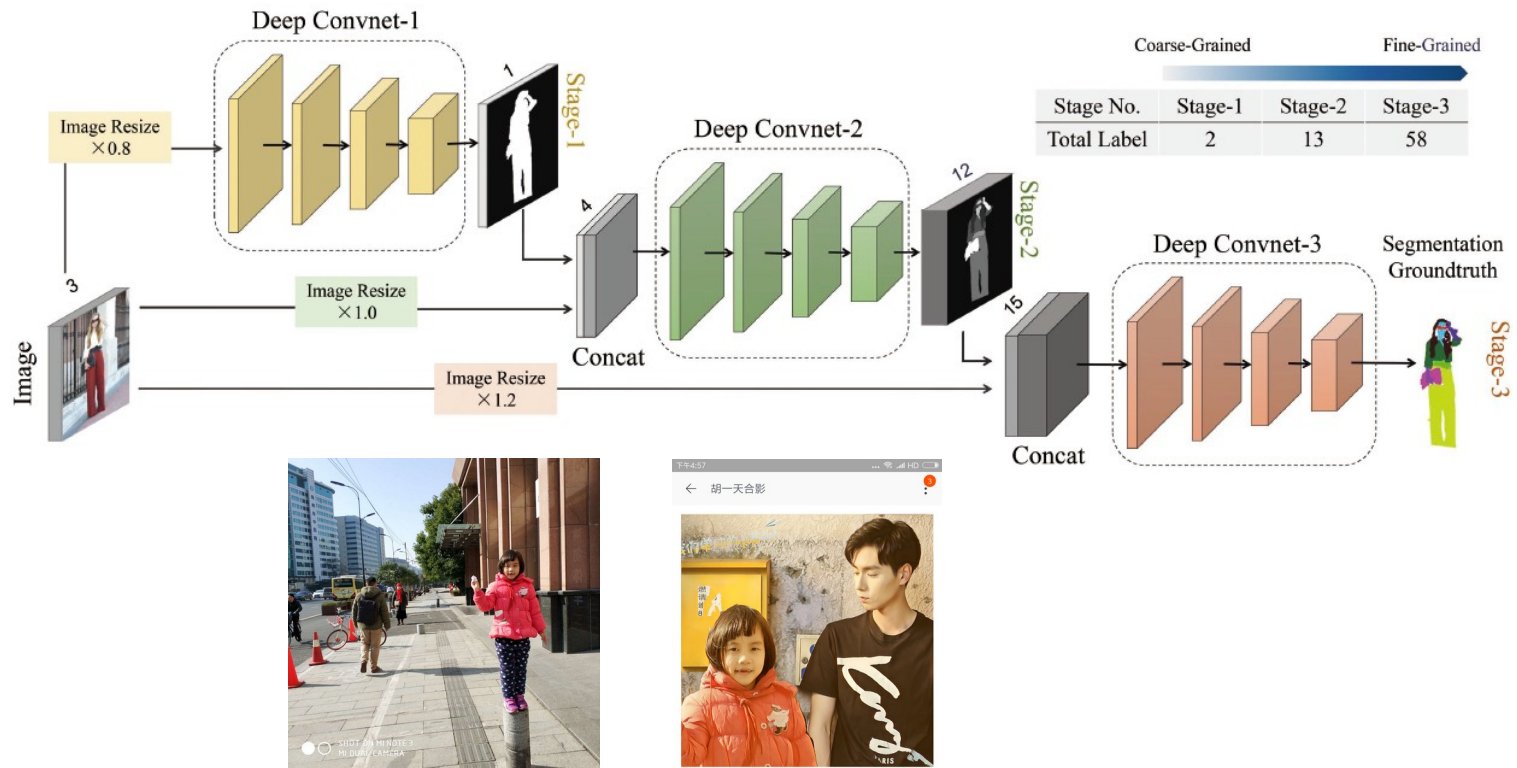
Logo detection




Semantic Structured Hashing for Person Re-identification



Finer-Net: Cascaded Human Parsing with Hierarchical Granularity



Neural Style Transfer




Input Content



Input Style

Neural Style Transfer



Output

R-CNN → **OverFeat** → MultiBox → SPP-Net → MR-CNN → DeepBox → AttentionNet →
2013.11 ICLR' 14 CVPR' 14 ECCV' 14 ICCV' 15 ICCV' 15 ICCV' 15

Fast R-CNN → DeepProposal → **Faster R-CNN** → **OHEM** → **YOLO v1** → G-CNN → AZNet →
ICCV' 15 ICCV' 15 NIPS' 15 CVPR' 16 CVPR' 16 CVPR' 16 CVPR' 16

Inside-OutsideNet(ION) → HyperNet → CRAFT → MultiPathNet(MPN) → **SSD** → GBDNet →
CVPR' 16 CVPR' 16 CVPR' 16 BMVC' 16 ECCV' 16 ECCV' 16


CPF → MS-CNN → **R-FCN** → PVANET → DeepID-Net → NoC → DSSD → TDM → **YOLO v2** →
ECCV' 16 ECCV' 16 NIPS' 16 NIPS' 16 PAMI' 16 TPAMI' 16 arXiv' 17 CVPR' 17 CVPR' 17

Feature Pyramid Net(**FPN**) → RON → DCN → DeNet → CoupleNet → **RetinaNet** → DSOD →
CVPR' 17 CVPR' 17 ICCV' 17 ICCV' 17 ICCV' 17 ICCV' 17 ICCV' 17

Mask R-CNN → SMN → **YOLO v3** → SIN → STDN → **RefineDet** → MLKP → Relation-Net →
ICCV' 17 ICCV' 17 arXiv' 18 CVPR' 18 CVPR' 18 CVPR' 18 CVPR' 18 CVPR' 18

Cascade R-CNN → RFBNet → CornetNet → Pelee → MethAnchor → SNIPER → **M2Det** ...
CVPR' 18 ECCV' 18 ECCV' 18 NIPS' 18 NIPS' 18 NIPS' 18 AAAI' 19

Training Neural Network is Difficult!



Remember:


Optimization through gradient descent

$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$$


How can we set the
learning rate?

Setting the Learning Rate

Small learning rate converges slowly and gets stuck in false local minima



Large learning rates overshoot, become unstable and diverge



How to deal with this?

Idea 1:

Try lots of different learning rates and see what works “just right”

How to deal with this?

Idea 1:

Try lots of different learning rates and see what works “just right”

Idea 2:

Do something smarter!

Design an adaptive learning rate that “adapts” to the landscape

Adaptive Learning Rates

- Learning rates are no longer fixed
- Can be made larger or smaller depending on:
 - how large gradient is
 - how fast learning is happening
 - size of particular weights
 - etc...

Adaptive Learning Rate Algorithms

- Momentum



`tf.train.MomentumOptimizer`

Qian et al. "On the momentum term in gradient descent learning algorithms." 1999.

- Adagrad



`tf.train.AdagradOptimizer`

Duchi et al. "Adaptive Subgradient Methods for Online Learning and Stochastic Optimization." 2011.

- Adadelata



`tf.train.AdadelataOptimizer`

Zeiler et al. "ADADELTA: An Adaptive Learning Rate Method." 2012.

- Adam



`tf.train.AdamOptimizer`

Kingma et al. "Adam: A Method for Stochastic Optimization." 2014.

- RMSProp




`tf.train.RMSPropOptimizer`

Additional details: <http://ruder.io/optimizing-gradient-descent/>

Gradient Descent

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights




Gradient Descent

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(W)}{\partial W}$
4. Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights


Can be very
computational to
compute!



Stochastic Gradient Descent

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
 3. Pick single data point i
 4. Compute gradient, $\frac{\partial J_i(\mathbf{W})}{\partial \mathbf{W}}$
 5. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
6. Return weights




Stochastic Gradient Descent

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick single data point i
4. Compute gradient, $\frac{\partial J_i(\mathbf{W})}{\partial \mathbf{W}}$
5. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
6. Return weights


Easy to compute but
very noisy
(stochastic)!



Stochastic Gradient Descent

Algorithm


1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick batch of B data points
4. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}} = \frac{1}{B} \sum_{k=1}^B \frac{\partial J_k(\mathbf{W})}{\partial \mathbf{W}}$
5. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
6. Return weights



Stochastic Gradient Descent

Algorithm


1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick batch of B data points
4. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}} = \frac{1}{B} \sum_{k=1}^B \frac{\partial J_k(\mathbf{W})}{\partial \mathbf{W}}$
5. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
6. Return weights




Fast to compute and a much better
estimate of the true gradient!

Overfitting


The Problem of Overfitting



Underfitting
Model does not have capacity
to fully learn the data




← Ideal fit →



Overfitting
Too complex, extra parameters,
does not generalize well

Regularization I: Dropout


- During training, randomly set some activations to 0



Regularization I: Dropout

- During training, randomly set some activations to 0
 - Typically 'drop' 50% of activations in layer
 - Forces network to not rely on any 1 node

 `tf.keras.layers.Dropout (p=0.5)`




Regularization I: Dropout

- During training, randomly set some activations to 0
 - Typically 'drop' 50% of activations in layer
 - Forces network to not rely on any 1 node




`tf.keras.layers.Dropout (p=0.5)`



Regularization 2: Early Stopping

- Stop training before we have a chance to overfit



Summary

- CNN is the best choice for many computer vision problems, but its limitations are also obvious:
 - Computational cost is very heavy (memory, time);
 - Supervised approach: needs lots of labeled data.
- Other options
 - Model distribution of input data instead of keeping their spatial information explicitly.
 - Unsupervised approach: can use unlabeled data
 - RBM, DBM, Deep Auto-Encoder, etc.